

▼ SGD Algorithm to predict movie ratings

There will be some functions that start with the word "grader" ex: `grader_matrix()`, `grader_mean()` those function definition.

Every Grader function has to return True.

1. Download the data from [here](#)
2. The data will be of this format, each data point is represented as a triplet of user_i

user_id	movie_id	rating
77	236	3
471	208	5
641	401	4
31	298	4
58	504	5
235	727	5

▼ Task 1

Predict the rating for a given (user_id, movie_id) pair


Predicted rating \hat{y}_{ij} for user i , movie j pair is calculated as $\hat{y}_{ij} = \mu + b_i + c_j + u_i^T v_j$, here we are using SGD algorithm with the optimization problem for N users and M movies is defined as

$$L = \min_{b, c, \{u_i\}_{i=1}^N, \{v_j\}_{j=1}^M} \alpha \left(\sum_j \sum_k v_{jk}^2 + \sum_i \sum_k u_{ik}^2 + \sum_i b_i^2 + \sum_{i,j \in \mathcal{I}^{\text{train}}} (y_{ij} - \mu - b_i - c_j - u_i^T v_j)^2 \right)$$

- μ : scalar mean rating
- b_i : scalar bias term for user i
- c_j : scalar bias term for movie j
- u_i : K-dimensional vector for user i
- v_j : K-dimensional vector for movie j

- *. We will be giving you some functions, please write code in that functions only.
- *. After every function, we will be giving you expected output, please make sure that you get that o

1. Construct adjacency matrix with the given data, assuming its [weighted un-directed bi-partite](#) rating given by user to the movie



Its Adjacency matrix

$$\begin{matrix}
 & \begin{matrix} 1 & 2 & 3 \end{matrix} \\
 \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 4 & 0 \\ 0 & 0 & 3 \\ 2 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 3 \end{bmatrix}
 \end{matrix}$$

you can construct this matrix like $A[i][j] = r_{ij}$ here i is user_id, j is movie_id and r_{ij} is rating giv

Hint : you can create adjacency matrix using [csr_matrix](#)

2. We will Apply SVD decomposition on the Adjacency matrix [link1](#), [link2](#) and get three matrices if A is of dimensions $N \times M$ then
 - U is of $N \times k$,
 - Σ is of $k \times k$ and
 - V is $M \times k$ dimensions.
 - *. So the matrix U can be represented as matrix representation of users, where each row u_i
 - *. So the matrix V can be represented as matrix representation of movies, where each row v movie.
3. Compute μ , μ represents the mean of all the rating given in the dataset. (write your code in `c`
4. For each unique user initialize a bias value B_i to zero, so if we have N users B will be a N di corresponds to the bias term for i^{th} user (write your code in `def initialize()`)
5. For each unique movie initialize a bias value C_j zero, so if we have M movies C will be a M will corresponds to the bias term for j^{th} movie (write your code in `def initialize()`)
6. Compute dL/db_i (Write you code in `def derivative_db()`)

7. Compute dL/dc_j (write your code in `def derivative_dc()`)

8. Print the mean squared error with predicted ratings.

for each epoch:

for each pair of (user, movie):

$b_i = b_i - \text{learning_rate} * dL/db_i$

$c_j = c_j - \text{learning_rate} * dL/dc_j$

predict the ratings with formula

$$\hat{y}_{ij} = \mu + b_i + c_j + \text{dot_product}(u_i, v_j)$$

9. you can choose any learning rate and regularization term in the range 10^{-3} to 10^2

10. **bonus:** instead of using SVD decomposition you can learn the vectors u_i, v_j with the help of

▼ Task 2

As we know U is the learned matrix of user vectors, with its i -th row as the vector u_i for user i . Each for a particular user.

The question we'd like to investigate is this: do our computed per-user features that are optimized to do with gender?

The provided data file [user_info.csv](#) contains an `is_male` column indicating which users in the data given the features U ?

Note 1 : there is no train test split in the data, the goal of this assignment is to give an intuition factorization with the help of SGD and application of truncated SVD. for better understanding check netflix case study.

Note 2 : Check if scaling of U, V matrices improve the metric

https://colab.research.google.com/drive/1_QHYQcHuKi227RUDzLIpu6U86npW8gKk#scrollTo=f8o7Z_JhaVbe&printMode=true



```
--2020-06-10 07:20:18-- https://doc-0k-0g-docs.googleusercontent.com/docs/securesc/45
Resolving doc-0k-0g-docs.googleusercontent.com (doc-0k-0g-docs.googleusercontent.com)...
Connecting to doc-0k-0g-docs.googleusercontent.com (doc-0k-0g-docs.googleusercontent.c
HTTP request sent, awaiting response... 200 OK
Length: 12073 (12K) [text/plain]
Saving to: 'user_info.csv.txt'
```

```
user_info.csv.txt 100%[=====>] 11.79K --.-KB/s in 0s
```

Reading the csv file

12/1591769175000/00484516897554883881/03543900857199698311/1-1z7iDB52cB6_Jp07Dqa-eOYSs-mivp

```
[-2020-06-10 06:06:55-- https://doc-0g-0g-docs.googleusercontent.com/docs/securesc/45
Resolving doc-0g-0g-docs.googleusercontent.com (doc-0g-0g-docs.googleusercontent.com)...
Connecting to doc-0g-0g-docs.googleusercontent.com (doc-0g-0g-docs.googleusercontent.c
HTTP request sent, awaiting response... 200 OK
Length: 880367 (860K) [text/csv]
Saving to: 'ratings_train.csv'
```

```
ratings_train.csv 100%[=====>] 859.73K --.-KB/s in 0.006s
```

```
2020-06-10 06:06:55 (130 MB/s) - 'ratings_train.csv' saved [880367/880367]
```

```
data1=pd.read_csv('user_info.csv.txt')
```

```
import pandas as pd
data=pd.read_csv('ratings_train.csv')
data.head()
```

```
[-]
  user_id  item_id  rating
0      772      36       3
1      471     228       5
2      641     401       4
3      312      98       4
4       58     504       5
```

```
data.shape
```

```
[-] (89992, 3)
```

Create your adjacency matrix

```
from scipy.sparse import csr_matrix
```

```
adjacency_matrix = csr_matrix((data.rating.values, (data.user_id.values, data.item_id.values
```

```
adjacency_matrix = csr_matrix((data.rating.values, (data.user_id.values, data.item_id.values
```

```
adjacency_matrix.shape
```

```
↳ (943, 1681)
```

Grader function - 1

```
def grader_matrix(matrix):
    assert(matrix.shape==(943,1681))
    return True
grader_matrix(adjacency_matrix)
```

```
↳ True
```

SVD decompostion

Sample code for SVD decompostion

```
from sklearn.utils.extmath import randomized_svd
import numpy as np
matrix = np.random.random((20, 10))
U, Sigma, VT = randomized_svd(matrix, n_components=5, n_iter=5, random_state=None)
print(U.shape)
print(Sigma.shape)
print(VT.T.shape)
```

```
↳ (20, 5)
    (5,)
    (10, 5)
```

Write your code for SVD decompostion

```
# Please use adjacency_matrix as matrix for SVD decompostion
# You can choose n_components as your choice
from sklearn.utils.extmath import randomized_svd
import numpy as np
matrix = np.random.random((943, 1681))
U, Sigma, VT = randomized_svd(adjacency_matrix, n_components=50, n_iter=50, random_state=None)
print(U.shape)
print(Sigma.shape)
print(VT.T.shape)
```

```
↳ (943, 50)
    (50,)
    (1681, 50)
```

Compute mean of ratings

```
def m_u(ratings):
    '''In this function, we will compute mean for all the ratings'''
    # you can use mean() function to do this
    # check this (https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.mean.html)

    return np.mean(ratings)
```

```
mu=m_u(data['rating'])
print(mu)
```

```
↳ 3.529480398257623
```

Grader function -2

```
def grader_mean(mu):
    assert(np.round(mu,3)==3.529)
    return True
mu=m_u(data['rating'])
grader_mean(mu)
```

```
↳ True
```

Initialize B_i and C_j

Hint : Number of rows of adjacent matrix corresponds to user dimensions(B_i), number of column dimensions (C_j)

```
def initialize(dim):
    '''In this function, we will initialize bias value 'B' and 'C'. '''
    # initialize the value to zeros
    # return output as a list of zeros
    temp=np.zeros(dim)

    return temp
```

```
dim=943 # give the number of dimensions for b_i (Here b_i corresponds to users)
b_i=initialize(dim)
```

```
dim=1681 # give the number of dimensions for c_j (Here c_j corresponds to movies)
c_j=initialize(dim)
```

Grader function -3

```
def grader_dim(b_i,c_j):
    assert(len(b_i)==943 and np.sum(b_i)==0)
    assert(len(c_j)==1681 and np.sum(c_j)==0)
    return True
grader_dim(b_i,c_j)
```

☞ True

Compute dL/db_i

```
def derivative_db(user_id,item_id,rating,U,V,mu,alpha):
    db=2*alpha*b_i[user_id]-2*(rating -mu-b_i[user_id]-c_j[item_id]- np.dot(U[user_id].T,V

    return db
```

Grader function -4

```
def grader_db(value):
    assert(np.round(value,3)==-0.931)
    return True
U1, Sigma, V1 = randomized_svd(adjacency_matrix, n_components=2,n_iter=5, random_state=24)
# Please don't change random state
# Here we are considering n_componets = 2 for our convinence
alpha=0.01
value=derivative_db(312,98,4,U1,V1,mu,alpha)
grader_db(value)
```

☞ True

Compute dL/dc_j

```
def derivative_dc(user_id,item_id,rating,U,V,mu):
    dc=2*alpha*c_j[item_id]-2*(rating -mu-b_i[user_id]-c_j[item_id]- np.dot(U[user_id].T,V

    return dc
```

Grader function - 5

```
def grader_dc(value):
    assert(np.round(value,3)==-2.929)
    return True
U1, Sigma, V1 = randomized_svd(adjacency_matrix, n_components=2,n_iter=5, random_state=24)
# Please don't change random state
# Here we are considering n_componets = 2 for our convinence
r=0.01
value=derivative_dc(58,504,5,U1,V1,mu)
```

```
grader_dc(value)
```

```
↳ True
```

Compute MSE (mean squared error) for predicted ratings

for each epoch, print the MSE value

for each epoch:

for each pair of (user, movie):

```
b_i = b_i - learning_rate * dL/db_i
```

```
c_j = c_j - learning_rate * dL/dc_j
```

predict the ratings with formula

$$\hat{y}_{ij} = \mu + b_i + c_j + \text{dot_product}(u_i, v_j)$$

```
from sklearn.metrics import mean_squared_error
```

```
mu=m_u(data["rating"])
```

```
y=data["rating"]
```

```
mes_y=[]
```

```
for epoch in range(0,50):
```

```
    y_i_j=[]
```

```
    for i in range(0,data.shape[0]):
```

```
        user_id1=data["user_id"][i]
```

```
        item_id1=data["item_id"][i]
```

```
        rating1=data['rating'][i]
```

```
        learing_rate=0.01
```

```
        b_i[user_id1]=b_i[user_id1]-learing_rate*(2*alpha*b_i[user_id1]-2*(rating1 -mu-b_i[use
```

```
        c_j[item_id1]=c_j[item_id1]-learing_rate*(2*alpha*c_j[item_id1]-2*(rating1 -mu-b_i[use
```

```
        y_i_j_temp=mu+b_i[user_id1]+c_j[item_id1]+np.dot(U[user_id1],VT[:,item_id1])
```

```
        y_i_j.append(y_i_j_temp)
```

```
    mes=mean_squared_error(y,y_i_j)
```

```
    print("for epoch = ",epoch,"mes = ",mes)
```

```
    mes_y.append(mes)
```

```
↳
```



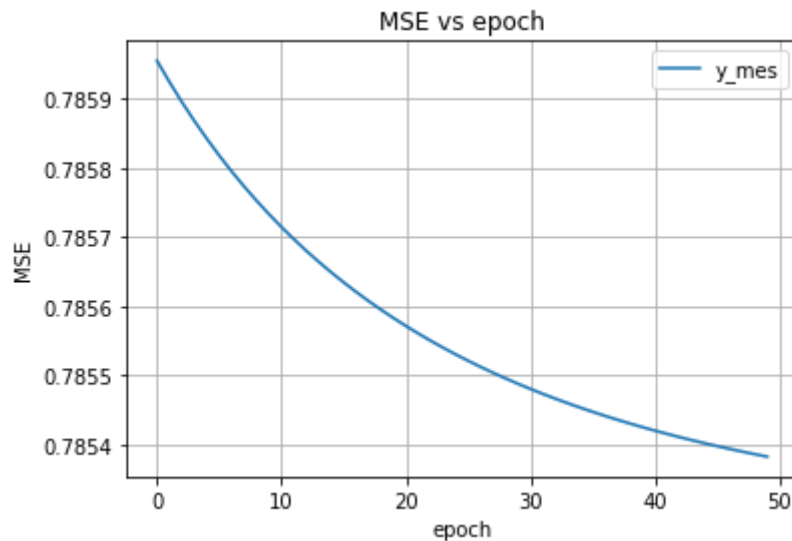
```
for epoch = 0 mes = 0.7859543895733097
for epoch = 1 mes = 0.7859236375626233
for epoch = 2 mes = 0.78589471555363
for epoch = 3 mes = 0.7858674751222213
for epoch = 4 mes = 0.7858417837327881
for epoch = 5 mes = 0.7858175226412596
for epoch = 6 mes = 0.7857945851210896
for epoch = 7 mes = 0.7857728749560172
for epoch = 8 mes = 0.7857523051542519
for epoch = 9 mes = 0.7857327968472996
for epoch = 10 mes = 0.78571427834341
for epoch = 11 mes = 0.7856966843110432
for epoch = 12 mes = 0.7856799550720919
for epoch = 13 mes = 0.7856640359880974
for epoch = 14 mes = 0.785648876925541
for epoch = 15 mes = 0.7856344317886024
for epoch = 16 mes = 0.7856206581096715
for epoch = 17 mes = 0.7856075166894547
for epoch = 18 mes = 0.7855949712797976
for epoch = 19 mes = 0.7855829883034146
for epoch = 20 mes = 0.7855715366055886
for epoch = 21 mes = 0.7855605872336466
for epoch = 22 mes = 0.7855501132406321
for epoch = 23 mes = 0.785540089510108
for epoch = 24 mes = 0.7855304925994631
for epoch = 25 mes = 0.785521300599456
for epoch = 26 mes = 0.7855124930080492
for epoch = 27 mes = 0.7855040506168415
for epoch = 28 mes = 0.7854959554086377
for epoch = 29 mes = 0.7854881904648812
for epoch = 30 mes = 0.785480739881845
for epoch = 31 mes = 0.785473588694609
for epoch = 32 mes = 0.78546672280798
for epoch = 33 mes = 0.7854601289336145
for epoch = 34 mes = 0.7854537945326822
for epoch = 35 mes = 0.7854477077635089
for epoch = 36 mes = 0.7854418574336799
for epoch = 37 mes = 0.7854362329561642
for epoch = 38 mes = 0.7854308243090566
for epoch = 39 mes = 0.785425621998589
for epoch = 40 mes = 0.7854206170250959
for epoch = 41 mes = 0.7854158008516555
for epoch = 42 mes = 0.7854111653751581
for epoch = 43 mes = 0.7854067028995784
for epoch = 44 mes = 0.7854024061112529
```

Plot epoch number vs MSE

- epoch number on X-axis
- MSE on Y-axis

```
#plot log_loss vs epoch
import matplotlib.pyplot as plt
a=list(range(0,50))
plt.plot(a,mes_y,label='y_mes')
plt.xlabel("epoch")
plt.ylabel("MSE")
plt.title("MSE vs epoch")
plt.legend()
```

```
plt.grid()
plt.show()
```



Task 2

```
data1=pd.read_csv('user_info.csv.txt')
data1.head()
```



	user_id	age	is_male	orig_user_id
0	0	24	1	1
1	1	53	0	2
2	2	23	1	3
3	3	24	1	4
4	4	33	0	5

▼ Training model logistic regression

As we know U is the learned matrix of user vectors, with its i -th row as the vector u_i for user i . Each for a particular user.

The question we'd like to investigate is this: do our computed per-user features that are optimized to do with gender?

The provided data file `user_info.csv` contains an `is_male` column indicating which users in the data given the features U ?

```
from sklearn import linear_model
```

```

Y=data1['is_male']
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
              early_stopping=False, epsilon=0.1, eta0=0.0001,
              fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
              loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
              penalty='l2', power_t=0.5, random_state=15, shuffle=True,
              tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)

clf = linear_model.SGDClassifier(alpha=0.0001,eta0=0.0001,max_iter=1000,loss='log',penalty
clf.fit(U,Y)

↳ SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                early_stopping=False, epsilon=0.1, eta0=0.0001,
                fit_intercept=True, l1_ratio=0.15, learning_rate='optimal',
                loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                penalty='l2', power_t=0.5, random_state=None, shuffle=True,
                tol=0.001, validation_fraction=0.1, verbose=0, warm_start=False)

user_male_pred = clf.predict(U)

```

▼ confusion_matrix

```

from sklearn.metrics import confusion_matrix
import seaborn as sns
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = (((C.T)/(C.sum(axis=1))).T)

    B =(C/C.sum(axis=0))
    plt.figure(figsize=(20,4))

    labels = [0,1]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")

    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')

```

```
plt.xlabel('Predicted Class')  
plt.ylabel('Original Class')  
plt.title("Recall matrix")
```

```
plt.show()
```

```
↳ /usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning  
import pandas.util.testing as tm
```

```
print('Train confusion_matrix')  
plot_confusion_matrix(Y,user_male_pred)
```

```
↳ Train confusion_matrix
```

